




Recurrent Neural Networks (RNNs)

Michel RIVEILL
michel.riveill@univ-cotedazur.fr



Order of words is important :



Motivation

- 3

What are RNNs?

- ▶ Main idea is to make use of sequential information
- ▶ How RNN is different from neural network?
 - ▶ Vanilla neural networks (MLP) **assume** all inputs are independent of each other
 - ▶ Features independence
 - ▶ But for many tasks, that's a very bad idea
- ▶ What RNN does?
 - ▶ Perform the same task for every element of a sequence
 - ▶ That's what **recurrent** stands for
 - ▶ Output depends on the previous computations!
- ▶ Another way of interpretation – RNNs have a “**memory**”
 - ▶ To store previous computations

Some applications (not recent)

▶  G  ED 

▶ - :/ - E3 | 

▶  G E  

▶  - :/ / / - 



▶  G  

▶ - :/ /2015/08/03/ - 


From vanilla NN to recurrent NN



▶ $y = F(U.X)$



From vanilla NN to recurrent NN



▶ $y = F(U.X)$




▶ $A \text{ } \boxed{\text{sequence of 10 squares}} h$



▶ $h_t = F(W \cdot h_{t-1} + U \cdot X_t)$ ✓✓


▶ $C \text{ } \boxed{\text{sequence of 10 squares}} h_t = F(V \cdot [h_{t-1}, X_t])$




From vanilla NN to recurrent NN



▶ $y = F(U.X)$



▶ $h_t = F(W.[h_{t-1}, X_t])$




▶ $h_1 = F(W.[h_0, X_1])$

▶ $h_2 = F(W.h_1, X_2]$

▶ $h_3 = F(W.h_2, X_3]$

▶ ...





From vanilla NN to recurrent NN



▶ $y = F(U.X)$



▶ $h_t = F(W.[h_{t-1}, X_t])$



"A"

"girl"

"walked"

...

"Certainly"



Remember

F

2



2

2



Remember





F




RNN in action



Problems with naive RNN

- ▶ 
- ▶  




Learning process

Vanishing gradient problem

<https://www.youtube.com/watch?v=8z3DFk4VxRo>



For rigorous proofs and derivations, please refer to

[On the difficulty of training recurrent neural networks, Pascanu *et al.*, 2013](#)
[Long Short-Term Memory, Hochreiter *et al.*, 1997](#)

Main learning problem



▶ $W_{new} = W_{old} - \lambda \text{ gradient}$



▶ $\frac{\partial E_3}{\partial W} = \sum_{k=0}^3 \frac{\partial E_3}{\partial \hat{y}_3} \frac{\partial \hat{y}_3}{\partial s_3} \frac{\partial s_3}{\partial s_k} \frac{\partial s_k}{\partial W}$

▶ $\frac{\partial E_3}{\partial W} = \frac{\partial E_3}{\partial \hat{y}_3} \frac{\partial \hat{y}_3}{\partial s_3} \frac{\partial s_3}{\partial W} + \frac{\partial E_3}{\partial \hat{y}_3} \frac{\partial \hat{y}_3}{\partial s_3} \frac{\partial s_3}{\partial s_2} \frac{\partial s_2}{\partial W} + \square + \frac{\partial E_3}{\partial \hat{y}_3} \frac{\partial \hat{y}_3}{\partial s_3} \frac{\partial s_3}{\partial s_2} \frac{\partial s_2}{\partial s_1} \frac{\partial s_1}{\partial s_0} \frac{\partial s_0}{\partial W}$

What is the value of the derivative 'chain'? $\frac{\partial s_3}{\partial s_2} \frac{\partial s_2}{\partial s_1} \frac{\partial s_1}{\partial s_0}$

► F

► = 0.25

► $0.25^2 = 0.0625$

► $0.25^4 = 0.00391$

► $0.25^8 = 0.0000152$

► $0.25^{16} = 0.0000000000233$

► 50


► G : 7 * 10^{-31}

►

► →

► :

Residual ?






How to introduce residual in RNN



From vanilla Sort Term Memory...

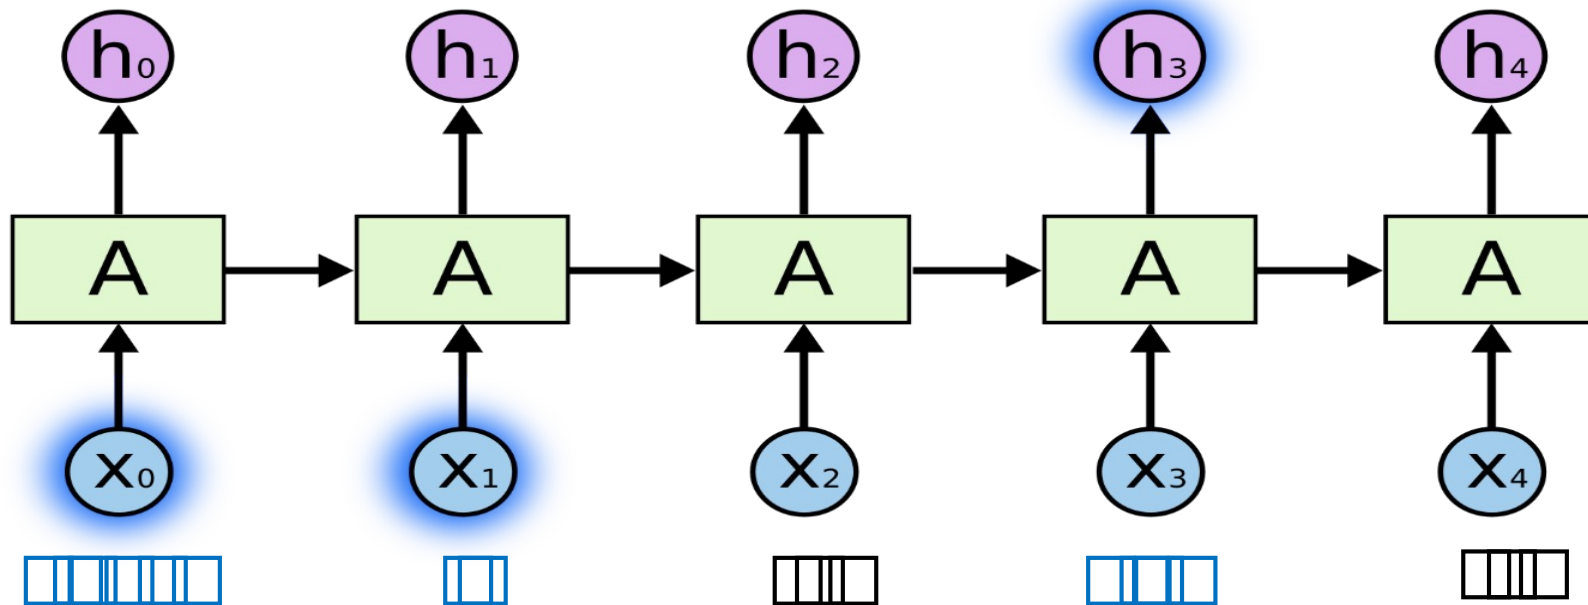


► F

►

► F


□







From vanilla Sort Term Memory...




▶ $\boxed{f_t = \tanh(W \boxed{h_{t-1}}, x_t)}$



... to LSTM (Long Short Term Memory)

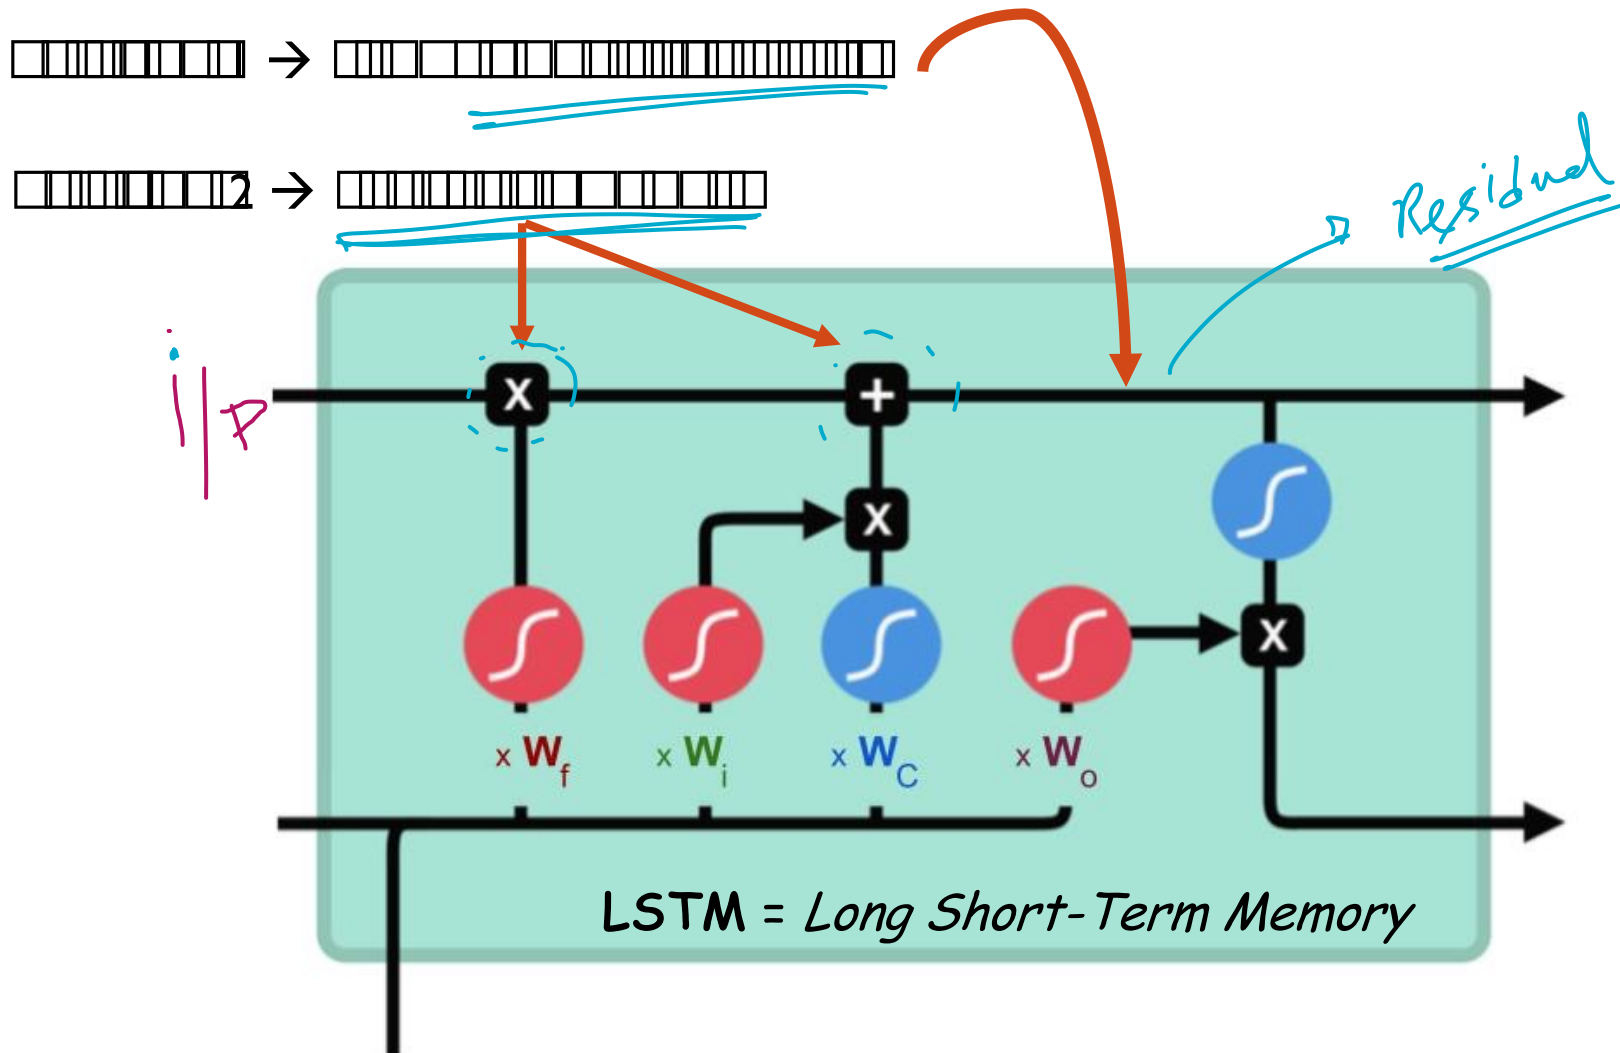
- ▶ B 
- ▶ 
- ▶ 
- ▶ 



|  ..




|  ..

Dealing with the vanishing gradient problem → LSTM cell



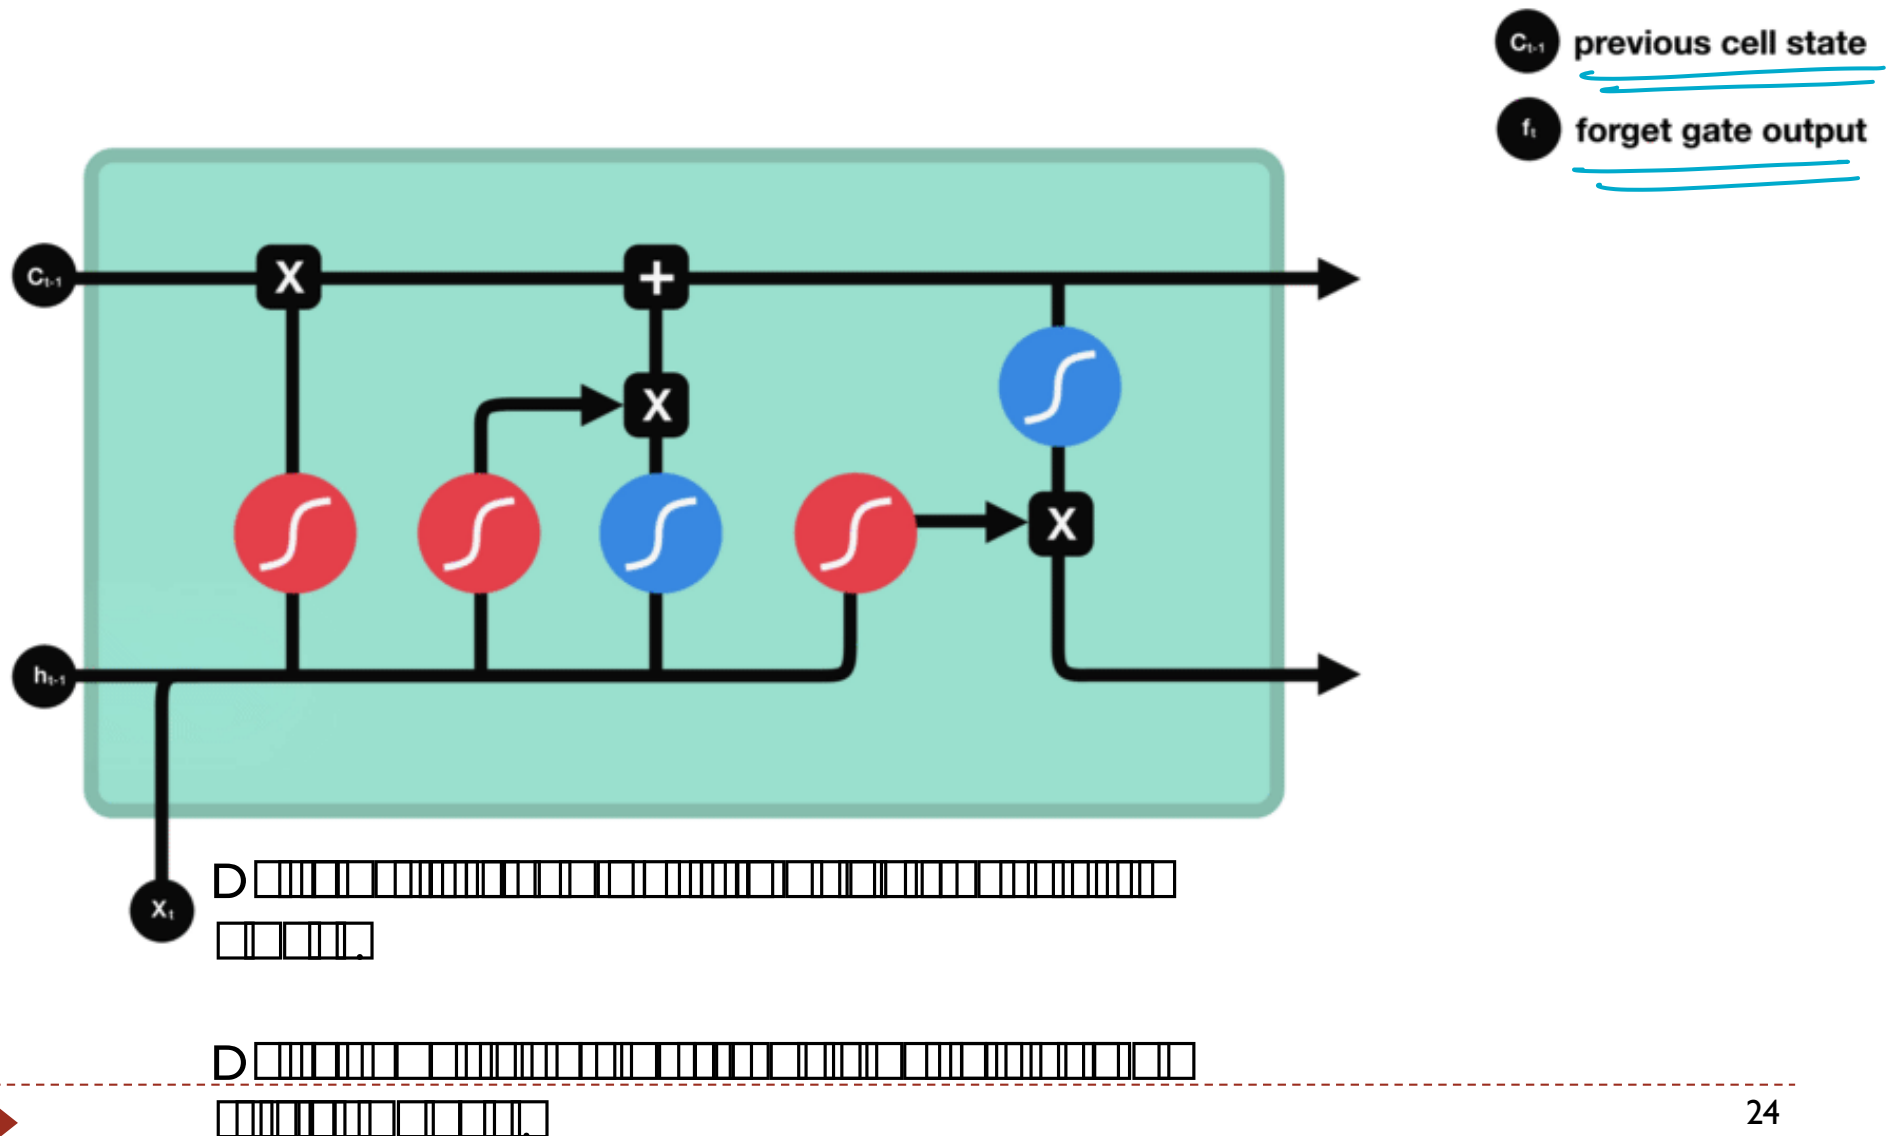
(crédit : image modifiée de Michaël Nguyen)

LSTM cell

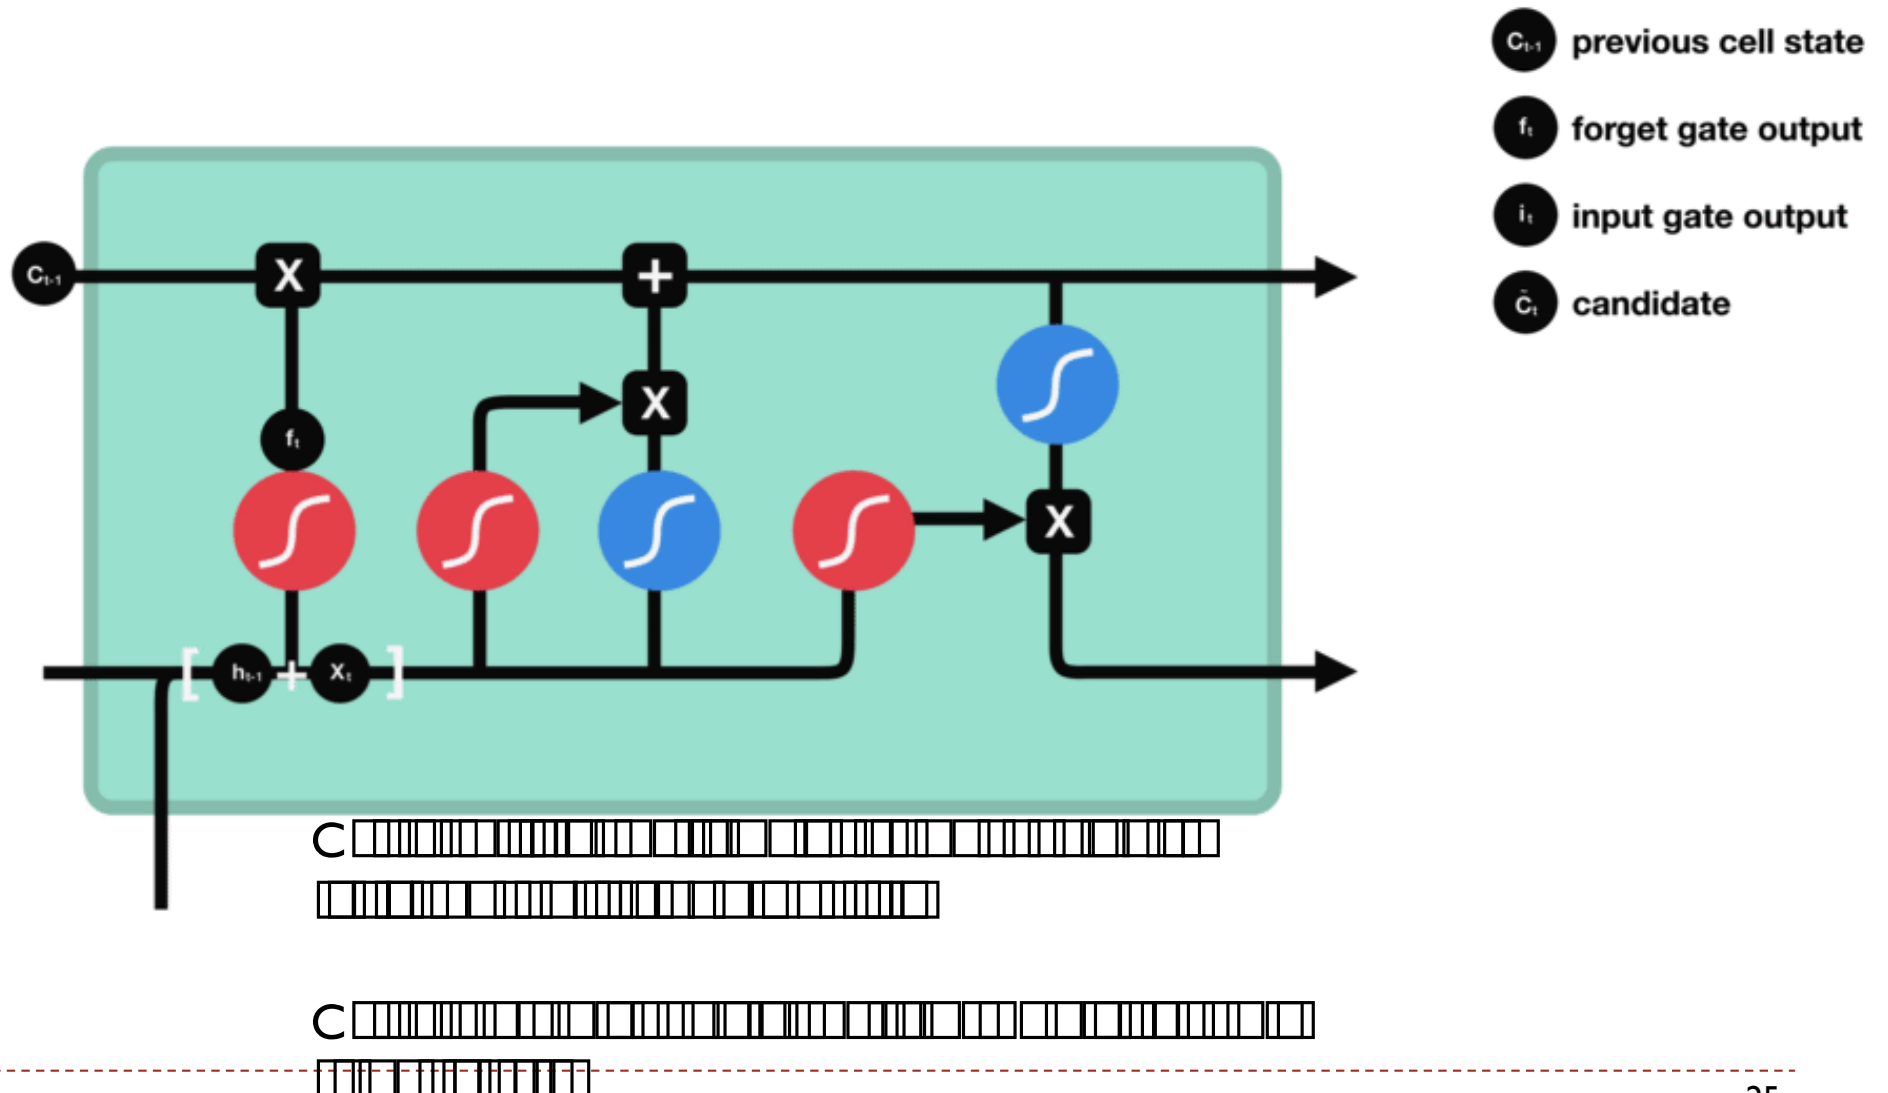


(crédit : image modifiée de Michaël Nguyen)



LSTM cell (porte oubli / forget get)




LSTM cell (porte entrée / input get)



LSTM cell (état de la cellule / cell state)











LSTM cell (porte de sortie / output gate)




La porte de sortie décide quel sera le prochain état caché. Il contient des informations sur les entrées précédentes du réseau et sert aux prédictions.

The output gate decides what the next hidden state will be. It contains information about previous inputs to the network and is used for predictions.

Naïve RNN vs LSTM

- ▶ 
- ▶ 
- ▶ 
- ▶ A  B 
, C 




c changes slowly



c^t is c^{t-1} added by something


h changes faster



h^t and h^{t-1} can be very different

GRU – gated recurrent unit

G = **C**



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$


$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$


- It combines the **forget** and **input** into a single **update gate**.
 - It also **merges** the **cell state** and **hidden state**.
- This is simpler/faster than LSTM.

Bi-directional RNNs


- ▶ RNNs can process the input sequence in forward and in the reverse direction




- Popular in speech recognition, could be used also with text




RNN cell in Keras




Keras Long Short-Term Memory Cell





`return_sequences = False`



`return_sequences = True`

**If there's a Dense after LSTM.
How many dense cells are used?**

- [illegible]

With  F ,

D  layer is applied only once at the last cell

With 

D  layer is applied to every timestep

↓
No of
Monomers

LS km \rightarrow 3D data (, ,)

A basic example: forecasting

$\text{forecast} = \text{model}(\text{features}(\text{data}, \text{index})) \# \text{features} = \text{data}[\text{index}]$

► $\text{features} = \text{data}, \text{index} = 5$ → One Feature.

$\text{forecast} = \text{model}(16, \text{features}(\text{data}, \text{index}))$

► $\text{features} = \text{data}, 16$

$\text{forecast} = \text{D}(\text{index}, \text{features}(\text{data}, \text{index}))$

► $\text{features} = \text{data}, 1$ → Regression.

F $\text{features}(\text{data}, \text{index})$

► $\text{features}(\text{data}, \text{index}) \leftarrow \text{features}(\text{data}, \text{index})$

F $\text{features}(\text{data}, \text{index})$

► $\text{features}(\text{data}, \text{index}) \leftarrow \text{features}(\text{data}, \text{index})$


► $\text{features}(\text{data}, \text{index})$



RNN for forecasting



RNN for forecasting

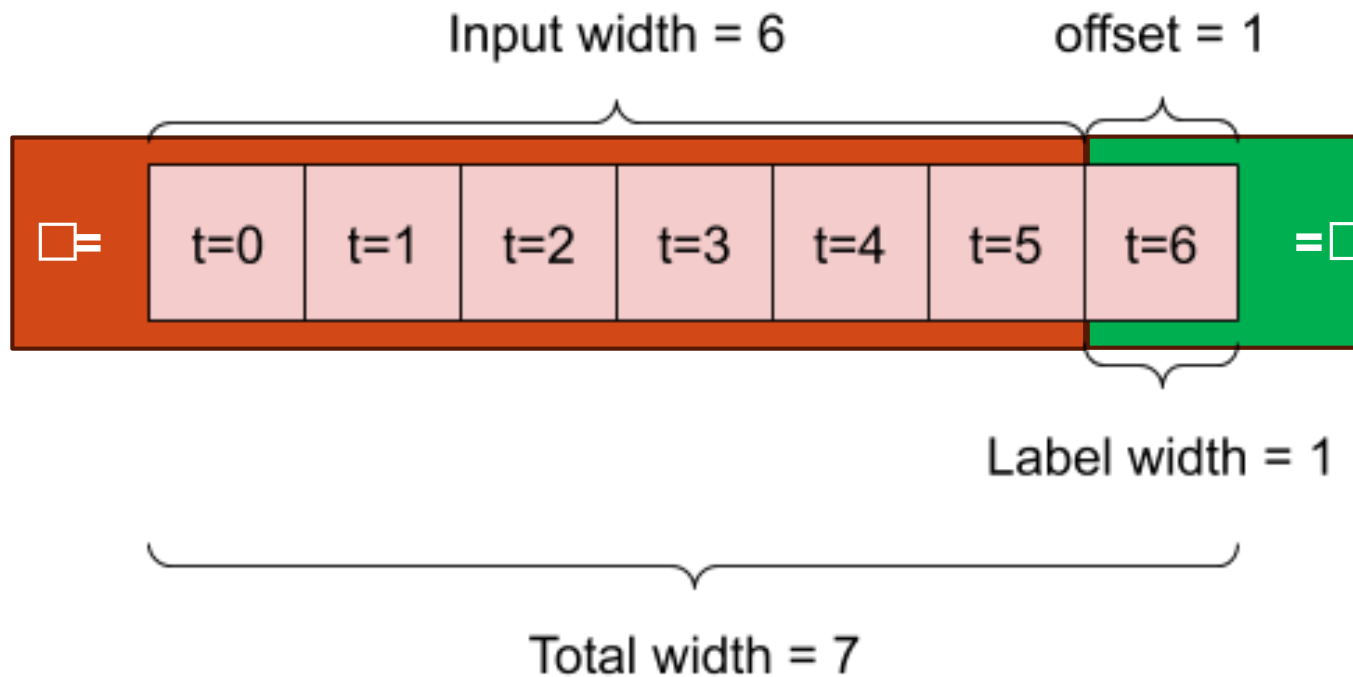


RNN for forecasting

▶ A 

▶ 


▶ 





▶ 24

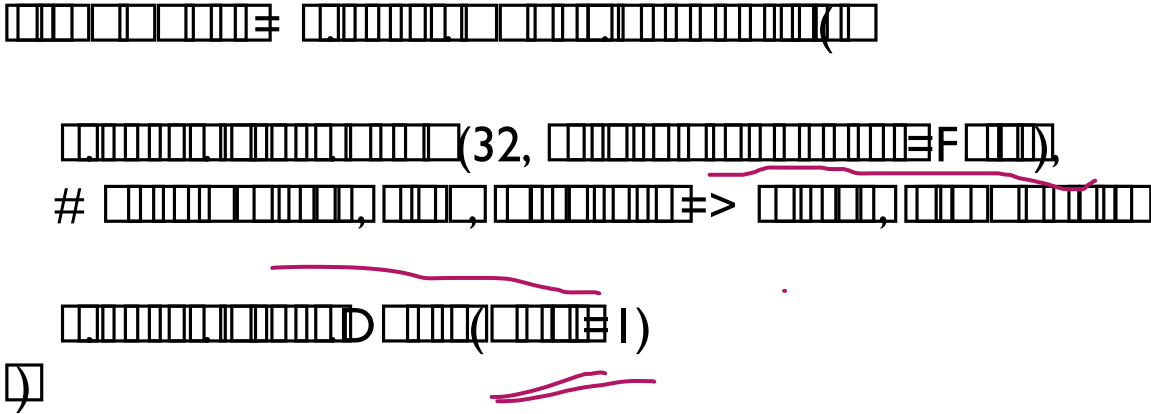
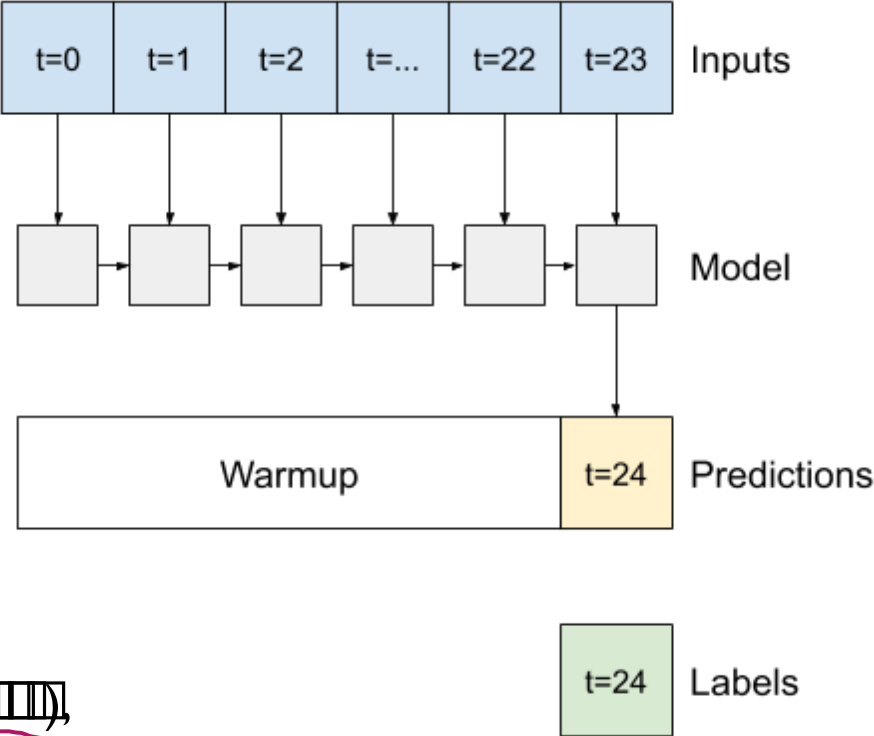
▶ 24



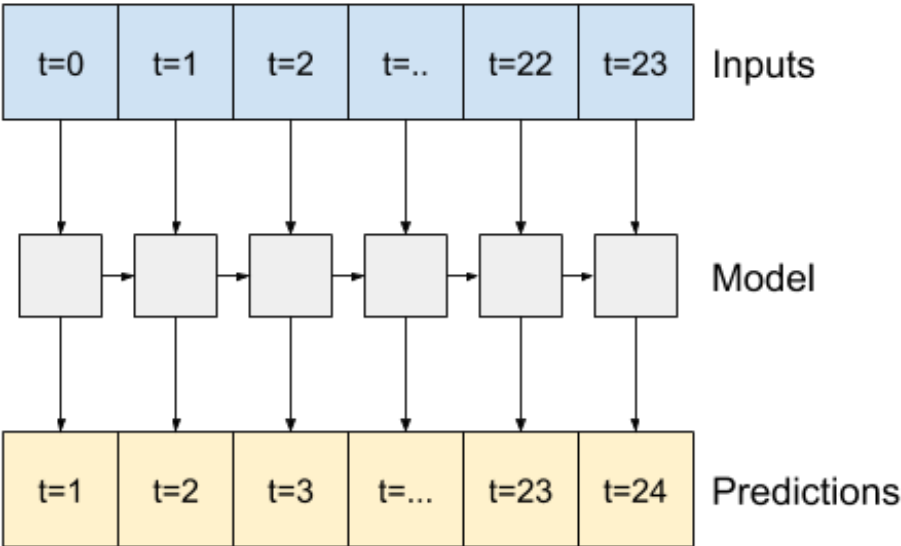
Dense model



Recurrent model: return_state=False



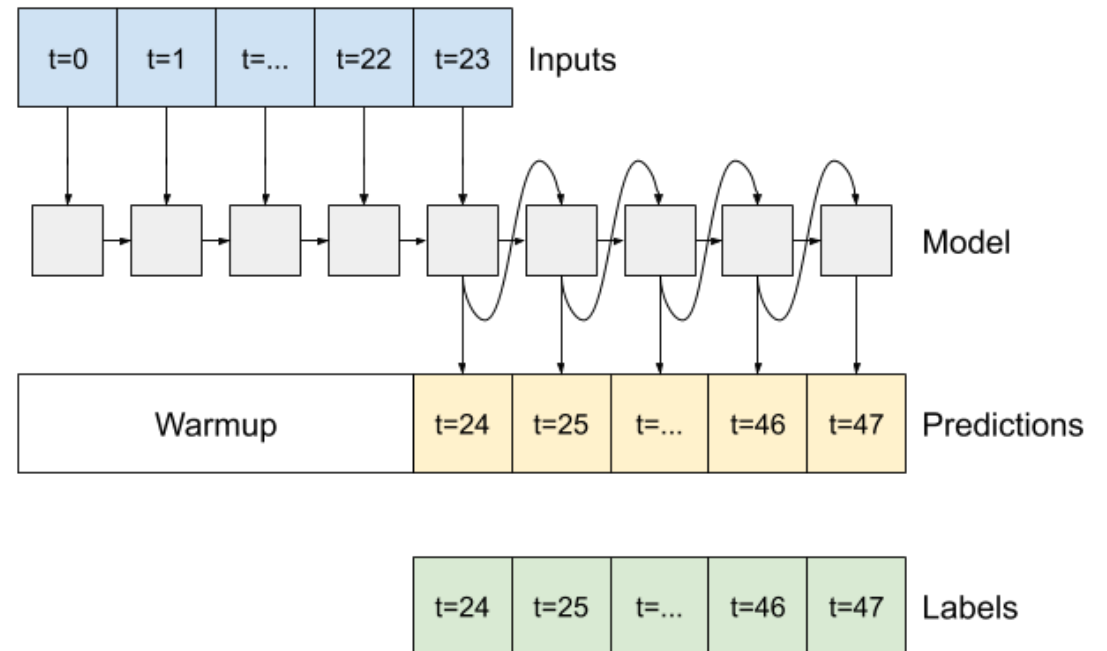
Recurrent model: return_state=True



), and a pink bracket underneath the 32 and the 3 elements."/>

 #


How to predict futur multiple values



Time series cross-validation



✓ `from sklearn.model_selection import TimeSeriesSplit`





RNN in NLP task






Some use of RNN

→ Text Classification / Sentiment analysis

A 

▶ C 

▶ 

▶  | DB



▶ I 

▶ 


▶ 

▶ 

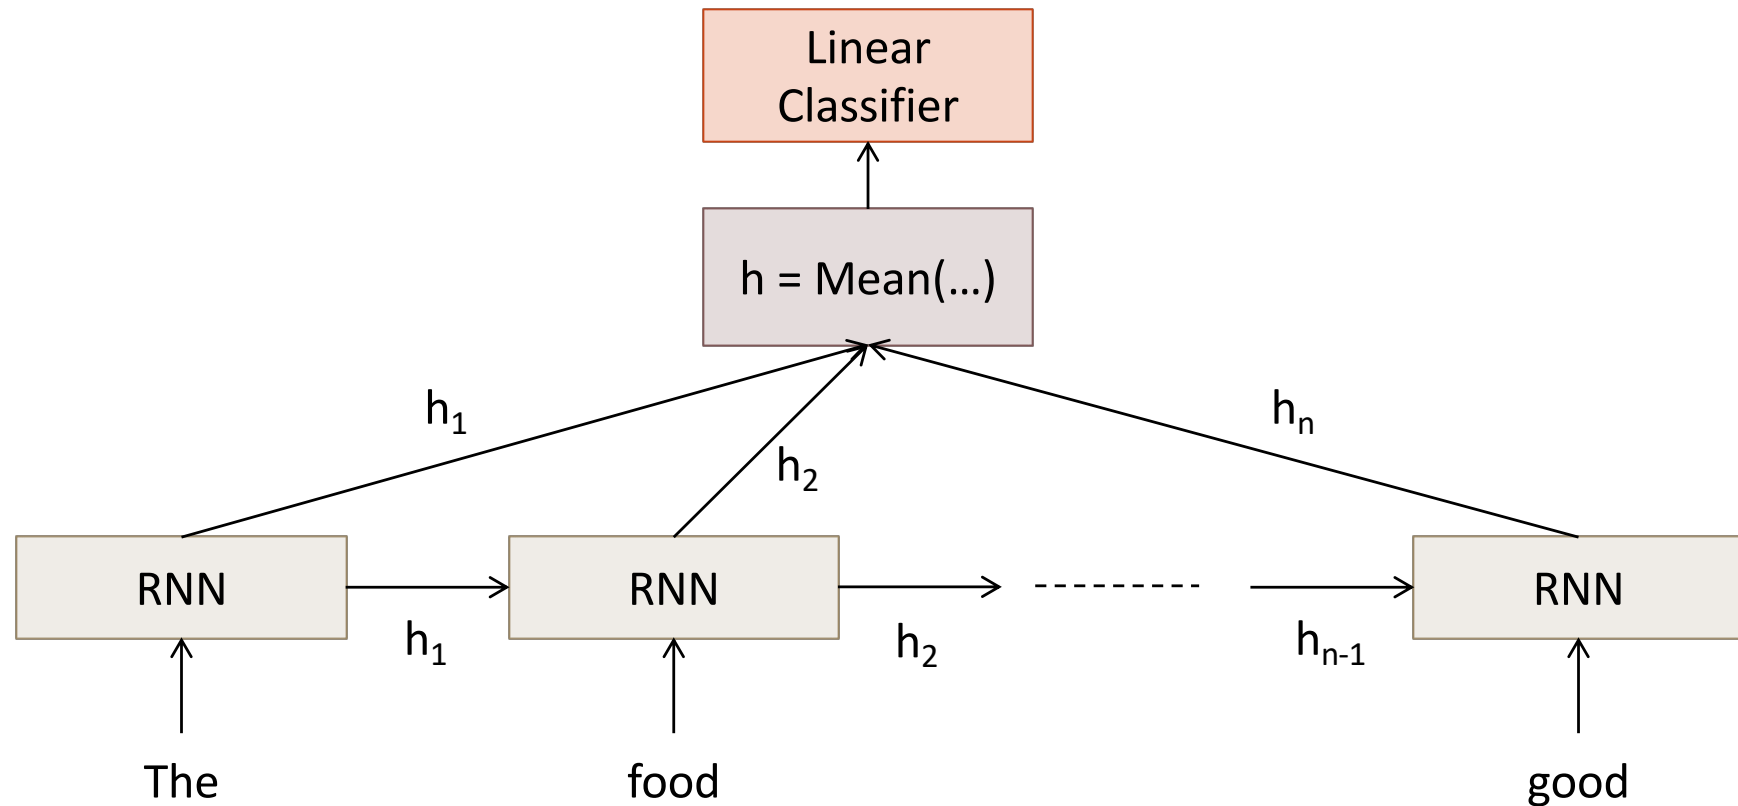
▶ 

▶ 

Sentiment analysis - solution 1



Sentiment analysis – solution2






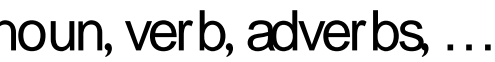

Some use of RNN

→ Named Entity Recognition / Part of Speech Tagging

▶ A 

▶ 

▶ C   () ✓

▶ C   () ✓

In fact, the **Chinese** **NORP** market has the **three** **CARDINAL** most influential names of the retail and tech space – **Alibaba** **GPE** , **Baidu** **ORG** , and **Tencent** **PERSON** (collectively touted as **BAT** **ORG**), and is betting big in the global **AI** **GPE** in retail industry space . The **three** **CARDINAL** giants which are claimed to have a cut-throat competition with the **U.S.** **GPE** (in terms of resources and capital) are positioning themselves to become the 'future **AI** **PERSON** platforms'. The trio is also expanding in other **Asian** **NORP** countries and investing heavily in the **U.S.** **GPE** based **AI** **GPE** startups to leverage the power of **AI** **GPE** . Backed by such powerful initiatives and presence of these conglomerates, the market in APAC AI is forecast to be the fastest-growing **one** **CARDINAL** , with an anticipated **CAGR** **PERSON** of **45%** **PERCENT** over **2018 - 2024** **DATE** .

To further elaborate on the geographical trends, **North America** **LOC** has procured **more than 50%** **PERCENT** of the global share in **2017** **DATE** and has been leading the regional landscape of **AI** **GPE** in the retail market. The **U.S.** **GPE** has a significant credit in the regional trends with **over 65%** **PERCENT** of investments (including M&As, private equity, and venture capital) in artificial intelligence technology. Additionally, the region is a huge hub for startups in tandem with the presence of tech titans, such as **Google** **ORG** , **IBM** **ORG** , and **Microsoft** **ORG** .

Extract information from text



Vente Villa 4 pièces Nice (06000)
Réf. 12390: Sur les Hauteurs de Nice. Superbe villa moderne (190m²), 2 chambres et 1 suite parentale, 3 salles de bain. Très grand salon/salle à manger, cuisine américaine équipée. Prestations de haut standing. Vue panoramique sur la mer. Cette villa a été construite en 2005. 1 270 000 euros. Si vous êtes intéressés, contactez vite Mimi LASOURIS 06.43.43.43. 43

REAL ESTATE TEMPLATE

Reference: 12390

Prize: 1 270 000

Surface: 190 m²

Year Built: 2005

Rooms: 4

Owner: Mimi LASOURIS

Telephone: 06.43.43.43. 43

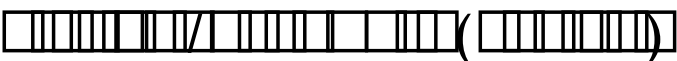
Mic

Machine Learning approach

► F 

► 

► C 


► 

► 


► 

► 


MLP for NER/POS



Recurrent neural network for NER

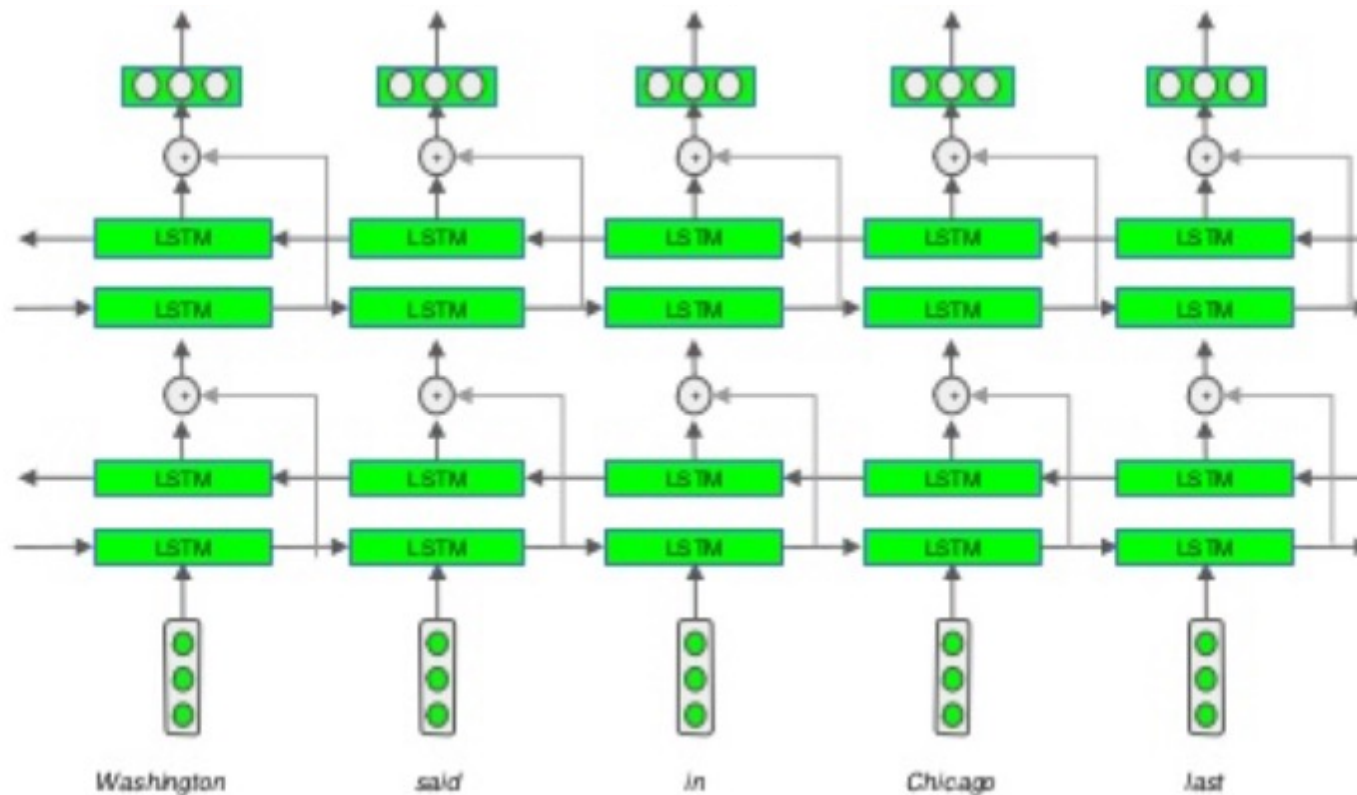




Recurrent neural network for NER (same network but unfolded)



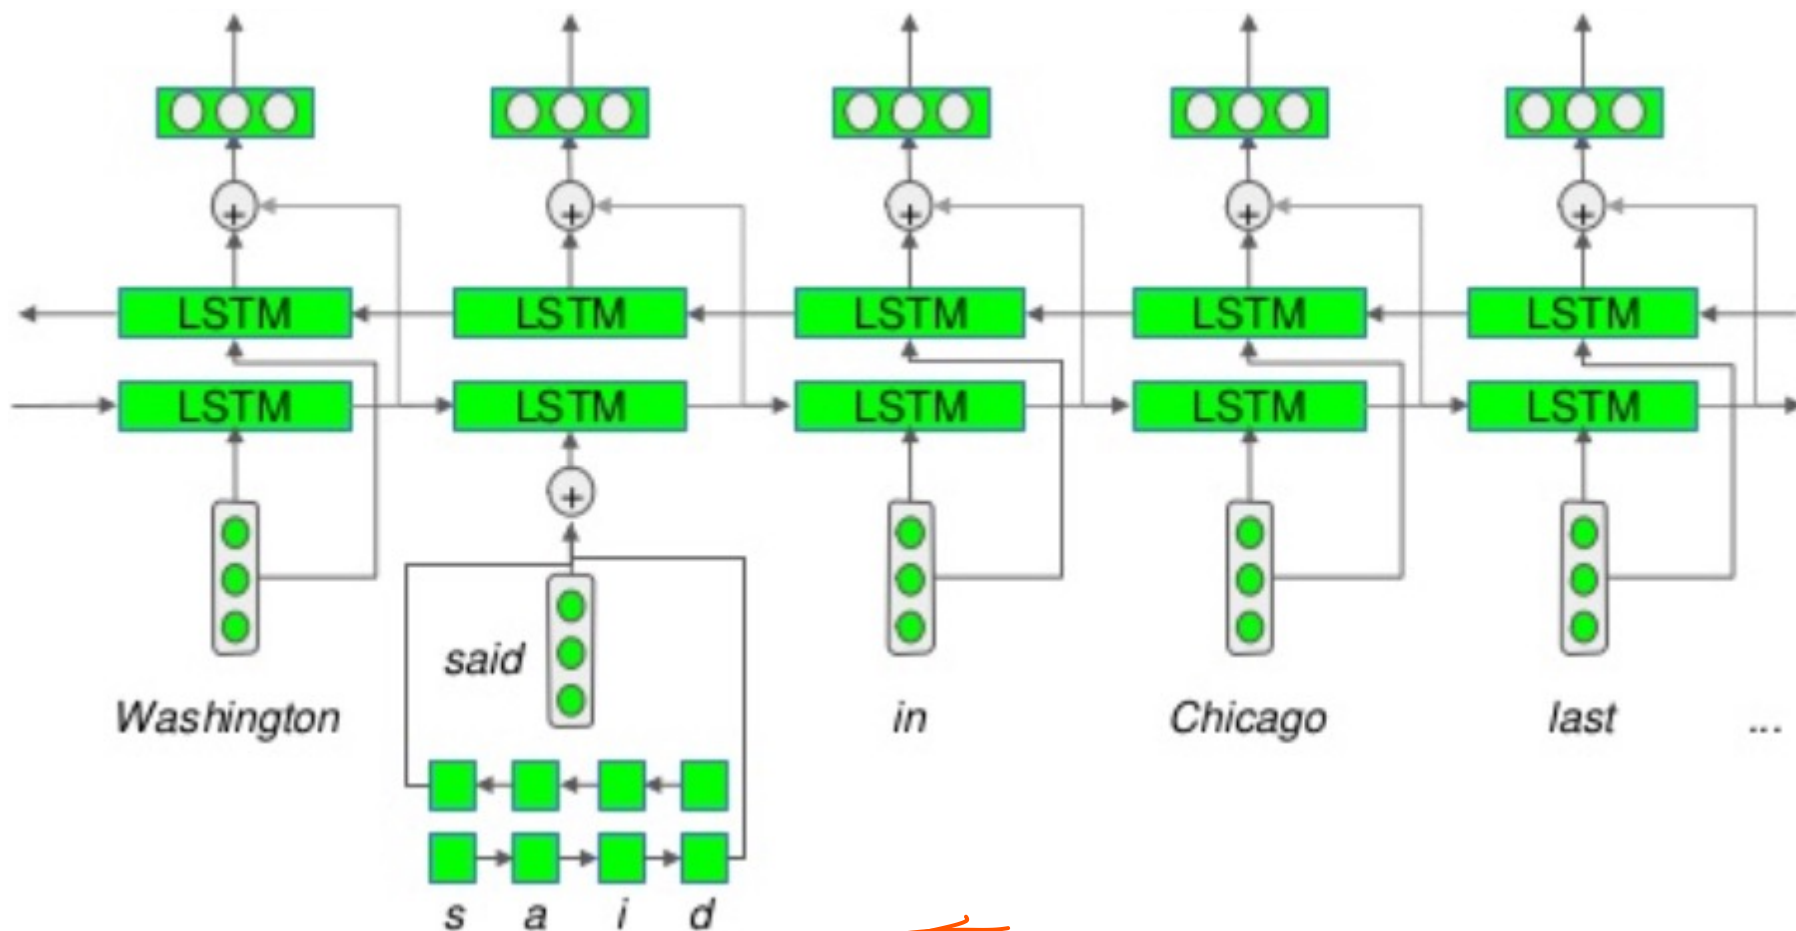
- A 

Stacked Bi-RNN









- A  

Multi-level encoding char encoding + word encoding



Today Lab

[illegible]

- **B** 
- 
- 
- 
- 
- 

▶ F

- [illegible]